

Detecting Roll Patterns in Copper Futures: A Statistical Analysis of Calendar Spread Dynamics

Austin Li

October 20, 2025

Abstract

This report presents an **hourly intraday-period analysis** of roll patterns in CME copper (HG) futures (2008–2024). We aggregate 1-minute data into **10 intraday periods per trading day** in US/Central time (7 one-hour US regular periods from 09:00–15:59, plus Late US, Asia, and Europe multi-hour periods). We detect **2,736 roll events** across **44,428 intraday periods** (6.2

For Framework Users

Unified CLI (daily/hourly) usage

1. Install dependencies: `pip install -r requirements.txt`
2. Hourly (intraday periods):
`python -m etf_roll_analysis.scripts.analyze --granularity hourly
--root ../organized_data/copper --settings
etf_roll_analysis/config/settings.yaml
--metadata etf_roll_analysis/metadata/contracts_metadata.csv
--output_dir etf_roll_analysis/outputs`
3. Daily:
`python -m etf_roll_analysis.scripts.analyze --granularity daily
--root ../organized_data/copper --settings
etf_roll_analysis/config/settings.yaml
--metadata etf_roll_analysis/metadata/contracts_metadata.csv
--output_dir etf_roll_analysis/outputs`

Full details in [Section 7: Framework Usage and Extension Guide](#).

Contents

1	Introduction	4
1.1	Background and Motivation	4
1.2	Research Objective	4
1.3	Dataset	4

2 Methodology	5
2.1 Data Processing Pipeline	5
2.1.1 Stage 1: Data Organization	5
2.1.2 Stage 2: Minute-to-Intraday-Period Aggregation	5
2.1.3 Stage 3: Panel Data Assembly	5
2.2 Front and Next Contract Identification	6
2.3 Calendar Spread Calculation	6
2.4 Roll Detection via Spread Widening	7
2.4.1 Rationale	7
2.4.2 Statistical Framework	7
2.4.3 Detection Criteria	7
2.4.4 Parameter Selection	8
2.4.5 Cool-down Mechanism	8
2.5 Timing Analysis	8
3 Results (Intraday Periods)	8
3.1 Intraday Distribution of Events	8
3.2 Hourly Timing (Days to Expiry)	9
3.3 Volume-Adjusted Preference Scores	9
3.4 Volume-Based Validation	10
4 Discussion	10
4.1 Interpretation of Results	10
4.1.1 Heterogeneous Roll Strategies	10
4.1.2 Absence of Universal Roll Date	11
4.1.3 Economic Implications	11
5 Project Structure and Documentation	11
5.1 Comprehensive Framework	11
5.2 Directory Organization	12
6 Framework Usage and Extension Guide	12
6.1 System Requirements and Installation	12
6.1.1 Prerequisites	12
6.1.2 Step-by-Step Installation	12
6.1.3 Troubleshooting Installation Issues	13
6.2 Understanding the Codebase Structure	13
6.2.1 Directory Organization	13
6.2.2 Data Flow Through the Pipeline	14
6.3 Running Your First Analysis	14
6.3.1 Default Analysis: Copper Futures	14
6.3.2 Understanding the Output Files	15
6.4 Extending to Other Commodities	15
6.4.1 Available Commodities	15
6.4.2 Switching to a Different Commodity	16
6.5 Customizing Analysis Parameters	17
6.5.1 Key Configuration Parameters	17
6.5.2 Parameter Impact on Results	17
6.6 Interpreting Analysis Results	18

6.6.1	Key Metrics to Examine	18
6.6.2	Statistical Summary	18
6.7	Common Issues and Solutions	19
6.7.1	Data-Related Issues	19
6.7.2	Configuration Issues	19
6.7.3	Performance Optimization	19
6.8	Advanced Usage	19
6.8.1	Batch Processing Multiple Commodities	19
6.8.2	Custom Visualizations	20
7	Conclusion	20
A	Quick Reference Card	23
A.1	Essential Commands	23
A.2	Key Configuration Parameters	23
A.3	Output Files Quick Guide	24
A.4	Commodity Folder Names	24
A.5	Key Statistical Formulas	24

1 Introduction

1.1 Background and Motivation

Futures contracts represent standardized agreements to buy or sell a specific quantity of a commodity at a predetermined price on a specified future date. Each contract has a defined expiration date, after which the contract must either be settled (physically or cash) or rolled into a subsequent contract month. For market participants maintaining continuous exposure—such as Exchange-Traded Funds (ETFs), commodity index funds, and hedge funds—the rolling process is essential for maintaining positions beyond individual contract lifespans.

The timing of roll activity has significant implications for market microstructure, price discovery, and trading strategies. Understanding when institutional participants systematically shift from front-month to next-month contracts can provide insights into:

- **Calendar spread dynamics:** Price relationships between adjacent contracts
- **Roll costs:** The transaction costs and slippage associated with position transfers
- **Market impact:** Price pressure created by large-scale rolling activity
- **Trading opportunities:** Potential arbitrage or momentum strategies

1.2 Research Objective

The primary objective of this analysis is to empirically determine *when* market participants roll their copper futures positions relative to contract expiration. Specifically, we aim to:

1. Detect roll events using statistical analysis of calendar spreads
2. Quantify the temporal distribution of rolls relative to expiry dates
3. Characterize the heterogeneity of roll strategies in the market
4. Develop a replicable framework applicable to other commodity futures

1.3 Dataset

Our analysis utilizes CME Group copper (ticker: HG) futures data with the following characteristics:

Parameter	Value
Commodity	Copper (HG)
Contracts	202 individual contracts
Date Range	January 1, 2008 – December 27, 2024
Trading Days	6,206
Data Frequency	1-minute bars
Data Aggregation	10 intraday periods per trading day (US/Central)
Analysis Periods	44,428 intraday periods
Contract Months	H, J, K, M, N, Q, U, V, X, Z (all 12 months traded)
Expiry Rule	Third last business day of contract month (CME)
Total Files Processed	13,548 futures files
Organized Structure	32 commodity folders in <code>organized_data/</code>
Data Inventory	<code>data_inventory.csv</code> mapping all files

Table 1: Dataset Characteristics

2 Methodology

2.1 Data Processing Pipeline

We aggregate 1-minute data to 10 intraday periods per trading day; identify front/next contracts using official CME expiry dates; compute calendar spreads as next minus front; detect widening via a z-score on consecutive intraday spread changes with a 20-period rolling window and a 3-hour time-based cool-down. Below is the high-level pipeline:

2.1.1 Stage 1: Data Organization

Input: 13,548 unorganized text files containing futures contract data across 32 commodity types.

Process: Automated categorization based on contract symbol pattern matching. Each filename follows the convention `SYMBOL_MONTHYEAR_1min.txt` (e.g., `HG_Z24_1min.txt`).

Output: Hierarchical directory structure with 32 commodity-specific folders. For copper, 202 individual contract files organized in `organized_data/copper/`.

2.1.2 Stage 2: Minute-to-Intraday-Period Aggregation

Input: 1-minute OHLCV data **Timezone:** US/Central (CME standard)

Aggregation: 10 intraday periods per day with hourly precision during US regular hours (periods 1–7) and multi-hour periods for Late US, Asia, and Europe (periods 8–10); period start timestamps are used for grouping.

Output: Intraday-period OHLCV per contract with period/session metadata.

2.1.3 Stage 3: Panel Data Assembly

Objective: Construct a wide-format panel where rows represent trading dates and columns represent individual contracts.

Structure:

- Dimensions: $6,206 \times 202$ (dates \times contracts)
- Each cell (t, c) contains the closing price of contract c on date t
- Metadata row: Official CME expiry dates for each contract
- Missing values (NaN) indicate contract not yet listed or already expired

Expiry Data Source: CME Group official futures calendar¹. All expiry dates calculated using CME rule: “Third last business day of the contract month.”

2.2 Front and Next Contract Identification

For each trading date t , we identify the active contracts:

Input: Panel data P , expiry dates E , current date t

Output: Front contract F_t , Next contract N_t

Active $\leftarrow \emptyset$;

foreach contract c in contracts **do**

if $P[t, c]$ is not NaN **and** $E[c] \geq t$ **then**
 | Add $(c, E[c])$ to Active;
 | **end**

end

Sort Active by expiry date (ascending);

if $|Active| = 0$ **then**

| $F_t \leftarrow \text{NaN}$, $N_t \leftarrow \text{NaN}$;

else

if $|Active| = 1$ **then**
 | $F_t \leftarrow \text{Active}[0]$, $N_t \leftarrow \text{NaN}$;
 | **else**
 | | $F_t \leftarrow \text{Active}[0]$, $N_t \leftarrow \text{Active}[1]$;
 | **end**

end

Algorithm 1: Front and Next Contract Identification

Coverage Results:

- Front contract identified: 5,284 days (85.1%)
- Next contract identified: 5,254 days (84.7%)
- Both front and next identified: 5,254 days (84.7%)

The 15% gap primarily occurs during early years (2008) when data coverage is incomplete.

2.3 Calendar Spread Calculation

The calendar spread S_t at time t is defined as:

$$S_t = P_{N_t}(t) - P_{F_t}(t) \quad (1)$$

¹<https://www.cmegroup.com/markets/metals/base/copper.calendar.html>

where $P_{N_t}(t)$ is the closing price of the next contract and $P_{F_t}(t)$ is the closing price of the front contract.

Interpretation:

- $S_t > 0$: Contango (next month premium)
- $S_t < 0$: Backwardation (next month discount)
- $\Delta S_t > 0$: Widening spread
- $\Delta S_t < 0$: Narrowing spread

Results: 5,254 valid spread observations (84.7% of trading days).

2.4 Roll Detection via Spread Widening

2.4.1 Rationale

When institutional participants systematically roll positions (selling front month, buying next month), this creates:

1. Downward pressure on front month prices
2. Upward pressure on next month prices
3. Net effect: Calendar spread widening

2.4.2 Statistical Framework

We employ a z-score methodology to detect statistically significant spread changes:

$$z_t = \frac{\Delta S_t - \mu_t}{\sigma_t} \quad (2)$$

where:

$$\Delta S_t = S_t - S_{t-1} \quad (\text{consecutive intraday-period spread change})$$

$$\mu_t = \frac{1}{w} \sum_{i=1}^w \Delta S_{t-i} \quad (\text{rolling mean, window } w)$$

$$\sigma_t = \sqrt{\frac{1}{w-1} \sum_{i=1}^w (\Delta S_{t-i} - \mu_t)^2} \quad (\text{rolling std dev})$$

2.4.3 Detection Criteria

A widening event W_t is flagged when:

$$W_t = \begin{cases} \text{True} & \text{if } z_t > \tau \text{ and } \Delta S_t > 0 \\ \text{False} & \text{otherwise} \end{cases} \quad (3)$$

2.4.4 Parameter Selection

Parameter	Value	Rationale
Window (w)	20 intraday periods	Short-term dynamics
Z-threshold (τ)	1.5	Event threshold for standardized changes
Min periods	10 intraday periods	Ensures statistical validity
Cool-down	3.0 hours	Time-based separation, more intuitive
Events detected	2,736	Total flagged during study window
Detection rate	6.2%	Share of intraday periods flagged

Table 2: Detection Parameters for Hourly Analysis

The choice of $\tau = 1.5$ standard deviations corresponds to approximately the 93rd percentile under normal distribution, flagging only substantial deviations from typical spread behavior.

2.4.5 Cool-down Mechanism

To prevent multiple signals for the same underlying roll event, we implement a 3-hour time-based cool-down at hourly granularity:

$$W_t = \begin{cases} \text{False} & \text{if } \exists t' \in [t - 3, t - 1] \text{ where } W_{t'} = \text{True} \\ \text{as above} & \text{otherwise} \end{cases} \quad (4)$$

2.5 Timing Analysis

For each detected widening event $W_t = \text{True}$, we calculate:

$$\text{Days to Expiry}_t = E[F_t] - t \quad (5)$$

where $E[F_t]$ is the expiration date of the front contract identified at time t .

3 Results (Intraday Periods)

3.1 Intraday Distribution of Events

The analysis detected **2,736 roll events** across 44,428 intraday periods (6.2% detection rate). The distribution reveals strong intraday patterns:

Intraday Period	Label	Events	Share
1	09:00 US Open	383	13.99%
2	10:00 US Morning	284	10.38%
3	11:00 US Late Morning	238	8.70%
4	12:00 US Midday	186	6.80%
5	13:00 US Early Afternoon	132	4.82%
6	14:00 US Late Afternoon	198	7.23%
7	15:00 US Close	151	5.52%
8	Late US/After-Hours (5h)	296	10.81%
9	Asia Session (6h)	454	16.59%
10	Europe Session (6h)	414	15.13%

Table 3: Distribution of roll events by intraday period

Key findings:

- **US Regular Hours:** 1,572 events (57.4% of total) - dramatic increase from 40.4%
- **Off-Peak Sessions:** 1,165 events (42.6% of total)
- **Hourly Density:** US hours average 224.6 events/hour vs 72.4 events/hour off-peak (3.1x higher)
- **Peak Single Hour:** 09:00 US Open with 383 events (14% of all events)

3.2 Hourly Timing (Days to Expiry)

Using hourly events aligned to front-contract expiries, days-to-expiry statistics: **median** 28 days, **mean** 30.7, IQR 16–43, range 0–116.

Count	Mean	Median	Std	P25	P75	Min	Max
2,736	30.65	28	19.32	16	43	0	116

Table 4: Days to expiry summary (hourly intraday-period events)

3.3 Volume-Adjusted Preference Scores

Preference scores adjust for volume patterns, with mean normalized to 1.0. Values above 1.0 indicate more roll activity than expected given volume share:

Rank	Intraday Period Label	Score
1	14:00 US Late Afternoon	2.22
2	15:00 US Close	1.89
3	09:00 US Open	1.05
4	13:00 US Early Afternoon	0.89
5	Late US/After-Hours	0.89

Table 5: Top 5 intraday periods by volume-adjusted preference

The end-of-day US hours (14:00-15:00) show **2x higher roll activity** than expected given their volume share, suggesting strategic timing around market close. Conversely, Asia (0.48) and Europe (0.52) show lower preference despite high absolute event counts, reflecting their high volume shares.

3.4 Volume-Based Validation

As an independent validation, we compute liquidity roll signals based on volume transitions:

$$L_t = \begin{cases} \text{True} & \text{if } V_{N_t}(t) \geq \alpha \cdot V_{F_t}(t) \\ \text{False} & \text{otherwise} \end{cases} \quad (6)$$

where $V_c(t)$ denotes volume of contract c at time t , and $\alpha = 0.8$ (80% threshold).

Results: 4,088 liquidity signals detected (65.9% of trading days).

Interpretation: The high frequency of liquidity signals relative to widening events (65.9% vs. 4.5%) suggests that:

1. Volume transitions are more gradual than discrete
2. Our widening detection captures only the most significant roll activity
3. Multiple smaller rolls may occur that don't trigger spread widening

4 Discussion

4.1 Interpretation of Results

4.1.1 Heterogeneous Roll Strategies

The wide distribution (IQR = 15 days, range = 66 days) and multiple modes indicate that copper futures market participants employ diverse roll strategies:

1. **Early Rollers (20–30 days before expiry):** Likely institutional players rolling large positions gradually to minimize market impact. The 21-day and 29-day modes suggest some participants follow fixed schedules.
2. **Middle Rollers (5–14 days before expiry):** The modal behavior, possibly representing the “optimal” trade-off between roll risk and timing costs. The 7-day and 12-day modes are prominent.
3. **Late Rollers (0–4 days before expiry):** Nearly 19% wait until the final days. These may include:
 - Hedgers with physical positions timed to delivery
 - Speculators attempting to profit from roll predictability
 - Smaller participants without liquidity concerns

4.1.2 Absence of Universal Roll Date

Unlike some ETF products (e.g., USO oil fund) that publicly disclose fixed roll schedules, the HG copper market exhibits no single dominant roll date. This suggests:

- No large single ETF dominates copper futures positioning
- Market participants strategically vary timing to reduce predictability
- Roll costs are spread across multiple days rather than concentrated

4.1.3 Economic Implications

For Roll Cost Analysis:

- Average roll window spans 15 days (IQR: 6–21)
- Roll premium/cost distributed rather than concentrated
- Early rolling may signal sophisticated market timing

For Trading Strategies:

- Calendar spread trades should focus on 5–21 day pre-expiry window
- No single “roll date” to front-run
- Volume analysis complements spread analysis for timing

5 Project Structure and Documentation

5.1 Comprehensive Framework

The analysis is now part of a fully documented Python framework with production-ready infrastructure:

- **Root-level README.md:** Complete project overview with installation instructions, quick start guide, and key findings
- **requirements.txt:** All Python dependencies (pandas, numpy, pyyaml, matplotlib, pytest)
- **setup.py:** Proper Python package structure for pip installation
- **Unit tests:** Test suite in `tests/` directory for validation
- **Relative paths:** All configurations use relative paths for portability across systems
- **Version control:** `.gitignore` configured for git repository

5.2 Directory Organization

```
futures_individual_contracts_1min/
|-- README.md                      # Main documentation
|-- requirements.txt                # Python dependencies
% USAGE.md removed in streamlined docs
|-- setup.py                        # Package installation
|-- organize_data.py               # Data organization script
|-- data_inventory.csv             # Maps 13,548 files
|-- organized_data/                 # 32 commodity folders
|   |-- copper/                     # 202 HG contracts
|   |-- gold/                       # GC contracts
|   |-- silver/                     # SI contracts
|   '-- [29 other commodities]
|-- etf_roll_analysis/              # Analysis framework
|   |-- config/settings.yaml        # Configuration
|   |-- src/roll_analysis/          # Core modules
|   |-- scripts/                    # Analysis scripts
|   '-- outputs/                   # Results
|-- tests/                          # Unit test suite
`-- presentation_docs/             # LaTeX reports
```

This structure enables researchers to easily extend the analysis to any commodity by simply changing one line in the configuration file.

6 Framework Usage and Extension Guide

This section provides comprehensive instructions for utilizing the futures roll analysis framework, designed to help users replicate our copper analysis and extend it to other commodities in the dataset.

6.1 System Requirements and Installation

6.1.1 Prerequisites

Before beginning, ensure your system meets the following requirements:

Component	Requirement
Python	Version 3.8 or higher
Operating System	Windows, Linux, macOS, or WSL
Memory	Minimum 4GB RAM (8GB recommended)
Disk Space	10GB for data + analysis outputs
LaTeX (optional)	For report compilation only

Table 6: System Requirements

6.1.2 Step-by-Step Installation

1. Navigate to the project directory:

```
cd futures_individual_contracts_1min/
```

2. Create a Python virtual environment (recommended):

```
# Windows
python -m venv venv
venv\Scripts\activate

# Linux/Mac/WSL
python3 -m venv venv
source venv/bin/activate
```

3. Install required packages:

```
pip install -r requirements.txt
```

4. Verify installation:

```
python -c "import pandas, numpy, yaml; print('All packages installed')"
```

5. (Optional) Install the package for development:

```
pip install -e .
```

6.1.3 Troubleshooting Installation Issues

- “**pip: command not found**”: Install pip with `python -m ensurepip`
- “**No module named yaml**”: The package name is `pyyaml`, not `yaml`
- **Permission errors**: Use `pip install --user` or activate a virtual environment
- **Version conflicts**: Create a fresh virtual environment to isolate dependencies

6.2 Understanding the Codebase Structure

6.2.1 Directory Organization

The framework follows a modular architecture designed for extensibility:

```
futures_individual_contracts_1min/
|-- organize_data.py           # Step 1: Organize raw data files
|-- organized_data/            # Result: 32 commodity folders
|   |-- copper/                # 202 HG contract files
|   |-- gold/                  # GC contracts
|   '-- [30 other commodities]
|-- etf_roll_analysis/        # Main analysis framework
|   |-- config/
|   |   '-- settings.yaml      # Configuration file
|   |-- src/roll_analysis/     # Core analysis modules
```

```

|   |   |-- ingest.py           # Data loading
|   |   |-- panel.py          # Panel assembly
|   |   |-- rolls.py          # Contract identification
|   |   |-- spread.py         # Spread calculation
|   |   '-- events.py        # Roll detection
|   |-- scripts/
|   |   '-- analyze.py       # Unified CLI (daily/hourly)
|   '-- outputs/
|   |   '-- panels/          # Price matrices
|   |   '-- roll_signals/    # Detected events
|-- presentation_docs/      # Reports and documentation

```

6.2.2 Data Flow Through the Pipeline

The analysis follows a sequential pipeline architecture:

1. **Data Ingestion** (`ingest.py`): Loads 1-minute CSV files, aggregates to intraday-period OHLCV (10 periods/day)
2. **Panel Assembly** (`panel.py`): Creates date \times contract price matrix
3. **Roll Identification** (`rolls.py`): Identifies front/next month contracts
4. **Spread Calculation** (`spread.py`): Computes calendar spreads
5. **Event Detection** (`events.py`): Applies z-score methodology
6. **Output Generation**: Saves results to CSV/Parquet files

6.3 Running Your First Analysis

6.3.1 Default Analysis: Copper Futures

The framework is pre-configured to analyze copper (HG) futures. To run the analysis:

1. Navigate to the analysis directory:

```

# Hourly (bucket)
python -m etf_roll_analysis.scripts.analyze --granularity
    hourly \
    --root ../organized_data/copper \
    --settings etf_roll_analysis/config/settings.yaml \
    --metadata etf_roll_analysis/metadata/contracts_metadata.
        csv \
    --output_dir etf_roll_analysis/outputs

# Daily
python -m etf_roll_analysis.scripts.analyze --granularity
    daily \
    --root ../organized_data/copper \
    --settings etf_roll_analysis/config/settings.yaml \
    --metadata etf_roll_analysis/metadata/contracts_metadata.
        csv \
    --output_dir etf_roll_analysis/outputs

```

2. Monitor the output:

```
Loading HG data from ../organized_data/copper
Found 202 contracts
Processing: HG_F08 , HG_G08 , HG_H08...
Building price panel...
Detecting roll events...
Found N widening events
Saving results to outputs/
Analysis complete!
```

6.3.2 Understanding the Output Files

The analysis generates four key output files:

File	Contents
panels/hg_panel_simple.csv	Price matrix (dates × contracts), 1.7MB
roll_signals/ hourly_widening.csv	Detected hourly roll events with timing
roll_signals/hg_spread.csv	Daily calendar spread time series
roll_signals/ hg_liquidity_roll.csv	Volume-based roll signals

Table 7: Output File Descriptions

6.4 Extending to Other Commodities

6.4.1 Available Commodities

The framework supports analysis of 32 commodity types:

Category	Commodity	Folder Name	Symbol
Metals	Copper	copper	HG
	Gold	gold	GC
	Silver	silver	SI
	Platinum	platinum	PL
	Palladium	palladium	PA
	Aluminum	aluminum	ALI
Energy	Crude Oil	crude_oil	CL
	Natural Gas	natural_gas	NG
	Heating Oil	heating_oil	HO
	Gasoline	gasoline	RB
	Brent Crude	brent_crude	BZ
Agriculture	Corn	corn	ZC
	Wheat	wheat	ZW
	Soybeans	soybeans	ZS
	Cotton	cotton	CT
	Sugar	sugar	SB
	Coffee	coffee	KC
	Cocoa	cocoa	CC
	Live Cattle	live_cattle	LE
Livestock	Lean Hogs	lean_hogs	HE
	Feeder Cattle	feeder_cattle	GF

Table 8: Available Commodities for Analysis

6.4.2 Switching to a Different Commodity

To analyze a different commodity, modify the configuration file:

1. Open the configuration file:

```
# From etf_roll_analysis directory
nano config/settings.yaml # or use any text editor
```

2. Change the data path:

```
# Original (copper):
data:
  minute_root: "../organized_data/copper"

# Change to gold:
data:
  minute_root: "../organized_data/gold"

# Or crude oil:
data:
  minute_root: "../organized_data/crude_oil"
```

3. Run the analysis:

```
python -m etf_roll_analysis.scripts.analyze --granularity
    hourly --root ../organized_data/gold
# or daily:
python -m etf_roll_analysis.scripts.analyze --granularity
    daily --root ../organized_data/gold
```

6.5 Customizing Analysis Parameters

6.5.1 Key Configuration Parameters

The config/settings.yaml file controls all analysis parameters:

Parameter	Default	Description
spread.window_buckets	20	Rolling window for z-score (in buckets)
spread.z_threshold	1.5	Standard deviations for event detection
spread.cool_down	3	Days between consecutive events
roll_rules.	0.8	Volume ratio for liquidity signals
liquidity_threshold		
data.price_field	close	Price field to use (close/settle)
data.timezone	US/Central	Market timezone

Table 9: Configuration Parameters

6.5.2 Parameter Impact on Results

Window Size (spread.window_buckets):

- **Smaller (10-15):** More responsive to recent changes, more signals
- **Default (20):** Balanced sensitivity and stability
- **Larger (25-30):** More stable, fewer false positives, may miss quick moves

Z-Score Threshold (spread.z_threshold):

- **Lower (1.0):** Detects 16% of events, more sensitive
- **Default (1.5):** Detects 5% of events, balanced
- **Higher (2.0):** Detects 2% of events, only extreme moves

Cool-down Period (spread.cool_down):

- **Shorter (1-2):** May capture multiple signals for same event
- **Default (3):** Prevents duplicate signals effectively
- **Longer (5-7):** May miss distinct but close events

6.6 Interpreting Analysis Results

6.6.1 Key Metrics to Examine

When reviewing the `hourly_widening.csv` output:

1. **days_to_expiry**: Most important metric - when rolls occur
2. **z_score**: Magnitude of the spread widening (higher = stronger signal)
3. **spread_change**: Actual dollar amount of spread widening
4. **front_contract**: Which contract is being rolled from
5. **next_contract**: Which contract is being rolled into

6.6.2 Statistical Summary

To generate summary statistics from the results:

```
import pandas as pd

# Load results
events = pd.read_csv('outputs/roll_signals/hourly_widening.csv')

# Summary statistics
print(f"Total events: {len(events)}")
print(f"Median days to expiry: {events['days_to_expiry'].median()}")
print(f"Mean days to expiry: {events['days_to_expiry'].mean():.2f}")
print(f"Std dev: {events['days_to_expiry'].std():.2f}")

# Distribution by range
ranges = [(0,4), (5,9), (10,14), (15,19), (20,24), (25,29),
          (30,100)]
for low, high in ranges:
    count = len(events[(events['days_to_expiry'] >= low) &
                        (events['days_to_expiry'] <= high)])
    pct = 100 * count / len(events)
    print(f"{low}-{high} days: {count} events ({pct:.1f}%)")
```

6.7 Common Issues and Solutions

6.7.1 Data-Related Issues

Issue	Solution
“No contracts found”	Check that data exists in the specified folder
“KeyError: ‘close’”	Verify CSV files have correct column names
“Memory error”	Process fewer contracts or increase system RAM
“NaN in spreads”	Normal for dates with missing data; framework handles it

Table 10: Common Data Issues

6.7.2 Configuration Issues

- **Path not found:** Use relative paths from `etf_roll_analysis` directory
- **YAML syntax error:** Check indentation (use spaces, not tabs)
- **Parameter out of range:** Window must be ≥ 2 , threshold ≥ 0

6.7.3 Performance Optimization

For large datasets or multiple commodities:

1. **Use Parquet format:** Faster than CSV for large panels
2. **Limit date range:** Add date filters in configuration
3. **Parallel processing:** Run multiple commodities simultaneously in separate terminals
4. **Memory management:** Process in chunks if memory limited

6.8 Advanced Usage

6.8.1 Batch Processing Multiple Commodities

Create a batch processing script:

```
# batch_analyze.py
import yaml
import subprocess
import os

commodities = ['copper', 'gold', 'silver', 'crude_oil', 'corn']

for commodity in commodities:
    # Update config
    with open('config/settings.yaml', 'r') as f:
        config = yaml.safe_load(f)
```

```

config['data']['minute_root'] = f"../organized_data/{commodity}"

with open('config/settings.yaml', 'w') as f:
    yaml.dump(config, f)

# Run analysis
print(f"Analyzing {commodity}...")
subprocess.run(['python', '-m', 'etf_roll_analysis.scripts.analyze', '--granularity', 'hourly'])

# Rename outputs
os.rename('outputs/roll_signals/hourly_widening.csv',
          f'outputs/roll_signals/{commodity}_hourly_widening.csv')

```

6.8.2 Custom Visualizations

Generate charts from the results:

```

import pandas as pd
import matplotlib.pyplot as plt

# Load results
events = pd.read_csv('outputs/roll_signals/hourly_widening.csv')

# Create histogram
plt.figure(figsize=(10, 6))
plt.hist(events['days_to_expiry'], bins=30, edgecolor='black')
plt.xlabel('Days to Expiry')
plt.ylabel('Number of Roll Events')
plt.title('Distribution of Roll Timing in Copper Futures')
plt.axvline(events['days_to_expiry'].median(), color='red',
            linestyle='--', label=f'Median: {events["days_to_expiry"].median():.0f} days')
plt.legend()
plt.grid(True, alpha=0.3)
plt.savefig('roll_distribution.png', dpi=300)
plt.show()

```

7 Conclusion

This hourly intraday-period analysis detected and characterized **2,736 roll events** in copper futures using a variable-granularity intraday structure. Key points:

- Intraday Pattern Discovery:** US regular hours show 3x higher roll density (224.6 events/hour) compared to off-peak sessions (72.4 events/hour)

2. **End-of-Day Concentration:** 14:00-15:00 CT exhibits 2x higher preference scores (2.22 and 1.89), indicating strategic timing around market close
3. **Detection summary:** 2,736 events over 44,428 intraday periods (6.2% of periods) using a 20-period window and 3-hour cool-down
4. **Intraday structure:** 7 one-hour US regular periods + 3 multi-hour session periods (Late US, Asia, Europe)
5. **Reproducible Framework:** Production-ready implementation for 32 commodity types with comprehensive documentation

The absence of a universal roll date in HG copper contrasts with some commodity ETFs that follow fixed schedules, suggesting a more sophisticated and distributed market microstructure. This has implications for:

- Calendar spread trading strategies
- Roll cost estimation for institutional portfolios
- Market making during roll periods
- Understanding copper futures market efficiency

The developed framework is modular and extensible, ready for application to the 31 additional commodity types in our organized dataset. The analysis framework has been packaged for academic research use. With installation guides (`requirements.txt`, `setup.py`) and test coverage (`tests/`), the framework is ready for extension to all 32 commodity types. The modular architecture and relative path configuration ensure portability and reproducibility across different computing environments.

Technical Implementation

Framework Architecture

The analysis pipeline consists of streamlined components:

- **Data Preparation:** `organize_data.py` organizes 13,548 files into 32 commodity folders
- **Core Modules** (in `src/roll_analysis/`):
 - `ingest.py`: Data loading and intraday-period aggregation
 - `bucket.py`: Intraday-period definitions (10 periods/day)
 - `bucket_panel.py`: Intraday-period panel assembly
 - `bucket_events.py`: Intraday-period event detection
 - `rolls.py`: Front/next contract identification
 - `spread.py`: Calendar spread calculation
 - `data_quality.py`: Contract filtering and quality checks
- **Analysis Script:** `scripts/analyze.py` (unified CLI; hourly/daily)

Configuration

All parameters externalized in `config/settings.yaml` using relative paths for portability:

```
data:
  minute_root: "../organized_data/copper" # Relative path
  timezone: "US/Central"
  price_field: "close"

spread:
  method: "zscore"
  window_buckets: 20
  z_threshold: 1.5
  cool_down_hours: 3.0

roll_rules:
  liquidity_threshold: 0.8
  confirm_days: 1
```

Installation and Reproducibility

The framework is now a proper Python package with comprehensive documentation:

```
# Install dependencies
pip install -r requirements.txt

# Install package in development mode (optional)
pip install -e .

# Run tests to verify installation
python -m pytest tests/

# Analyze any commodity (module invocation)
# 1. Edit config/settings.yaml to change commodity path
# 2. Run analysis via unified CLI
python -m etf_roll_analysis.scripts.analyze --granularity hourly \
    \
    --root ../organized_data/copper \
    --settings etf_roll_analysis/config/settings.yaml \
    --metadata etf_roll_analysis/metadata/contracts_metadata.csv \
    --output_dir etf_roll_analysis/outputs
```

Complete documentation available:

- `README.md`: Project overview and quick start guide
- `requirements.txt`: All Python dependencies
- `setup.py`: Package installation configuration

All analysis outputs stored in `outputs/` directory:

- `panels/hg_panel_simple.csv`: Price panel (1.7 MB)

- `roll_signals/hg_widening.csv`: Detected events (281 rolls)
- `roll_signals/hg_spread.csv`: Calendar spreads time series
- `roll_signals/hg_liquidity_roll.csv`: Volume-based signals

References

1. CME Group. “Copper Futures Contract Specifications.” <https://www.cmegroup.com/markets/metals/base/copper.contractSpecs.html>
2. CME Group. “Copper Futures Calendar.” <https://www.cmegroup.com/markets/metals/base/copper.calendar.html>
3. CME Group Rulebook. “Chapter 112: Copper Futures.” Accessed 2024.

A Quick Reference Card

A.1 Essential Commands

Quick Command Reference

```
# Setup (one-time)
cd futures_individual_contracts_1min/
pip install -r requirements.txt
python organize_data.py # If data not organized

# Run Hourly Analysis
python -m etf_roll_analysis.scripts.analyze --granularity
    hourly \
    --root ../organized_data/copper \
    --settings etf_roll_analysis/config/settings.yaml \
    --metadata etf_roll_analysis/metadata/contracts_metadata.
        csv \
    --output_dir etf_roll_analysis/outputs

# Change Commodity
# Edit config/settings.yaml:
# data.minute_root: "../organized_data/gold"
```

A.2 Key Configuration Parameters

Parameter	Default	Effect
<code>spread.window_buckets</code>	20	Sensitivity window (2 days)
<code>spread.z_threshold</code>	1.5	Detection rate (6.16%)
<code>spread.cool_down_hours</code>	3.0	Time-based event separation
<code>bucket_config.enabled</code>	true	Enable hourly analysis

A.3 Output Files Quick Guide

File	Contents
panels/hg_panel_simple.csv	Price matrix
roll_signals/ hourly_widening.csv	Hourly roll events
roll_signals/hg_spread.csv	Calendar spreads
roll_signals/ hg_liquidity_roll.csv	Volume signals

A.4 Commodity Folder Names

Metals: copper, gold, silver, platinum, palladium

Energy: crude_oil, natural_gas, heating_oil, gasoline

Agriculture: corn, wheat, soybeans, cotton, sugar, coffee

Livestock: live_cattle, lean_hogs, feeder_cattle

A.5 Key Statistical Formulas

$$\text{Z-score} = \frac{\Delta S_t - \mu_{20}}{\sigma_{20}}$$

$$\text{Calendar Spread} = P_{\text{next}} - P_{\text{front}}$$

$$\text{Detection} = z_t > 1.5 \text{ and } \Delta S_t > 0$$

Full documentation: See Section 7 of this report